



Binary Prediction of Poisonous Mushrooms

Dataset: Kaggle Playground Series (S4E8)

Course: CSE 572 Data Mining

Group Members: Luna Sbahtu, Molly Shircliff, Bright Mudzingwa, Pardon Hlongwane

Checkpoint: 2 – *Predictive Methods*



Checkpoint 2: Overview

- Predictive Methods
- Results & Issues
- Improvements Proposed & Results
- Future Steps



Predictive Methods

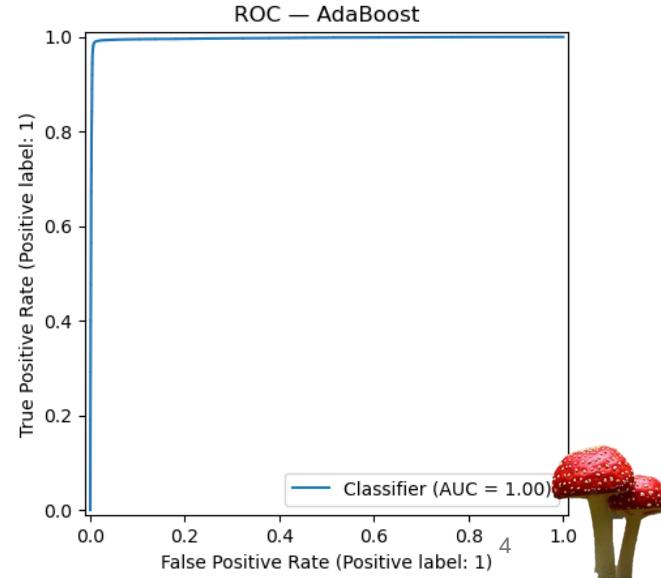
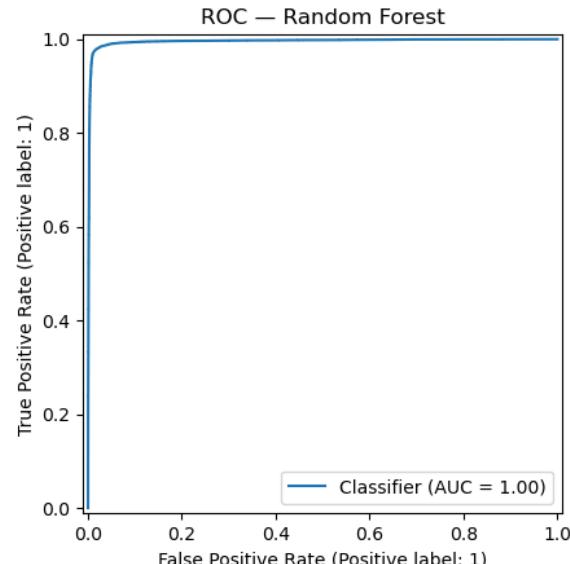
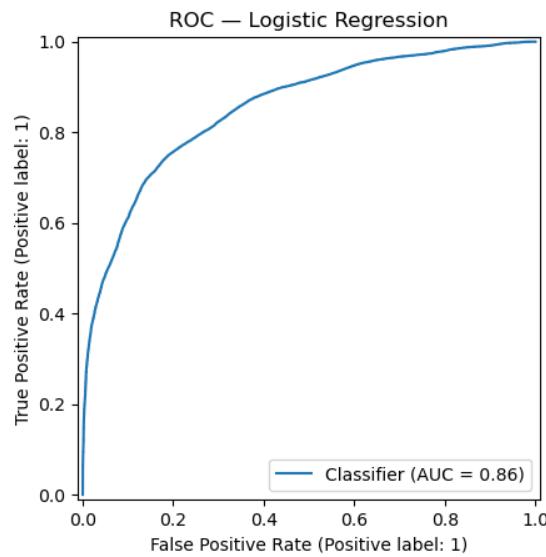
- Chose to perform 3 separate models:
 - Logistic Regression Classifier
 - Use this model as a baseline. Is there a linear relationship between classes?
 - Random Forest Classifier
 - This is preferred if there is nonlinear relationship. It is better than Decision Tree to check multiple split combinations
 - AdaBoost
 - This is a further improvement on random forest, but slower. Will it be worth the computational tradeoff?



Model Baseline Results

Due to the large dataset ($> 3\text{m}$ rows), models were created with 10% random sampling with replacement. In order to keep class proportions balanced, a stratified split was also applied when creating the train-test split for modeling.

Model	Accuracy	F1	ROC
Logistic Regression	78.02%	0.773	0.8551
Random Forest	98.06%	0.981	0.995
AdaBoost	98.99	0.9899	0.996



Baseline Model Discussion



Logistic Regression performed worst
(78.02% accuracy)

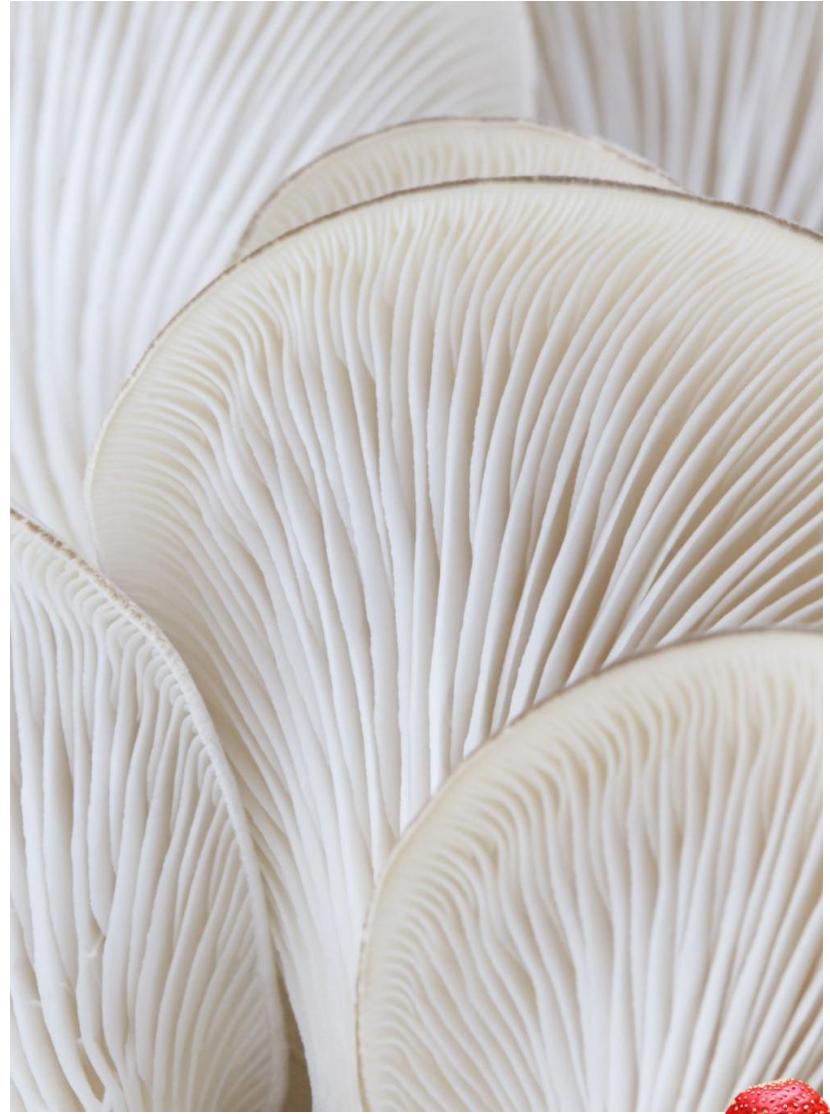
- Potential causes:
 - Non-linear boundary between edible & poisonous class
 - Feature interactions



Random Forest & AdaBoost performed best, with AdaBoost performing slightly better (0.9% improvement)

Decision:

- Choose Random Forest & AdaBoost to perform improvements on



Improvements Proposal

5-fold cross validation

Ensure that accuracy values are correct and not overfitted

Grid search & hyperparameters analysis

Determine best parameters for better error metrics



5-Fold Cross Validation

Model	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Avg
Random Forest	97.85	97.88	98.18	97.97	98.1	98.05%
AdaBoost	98.90	98.94	98.94	98.95	98.94	98.93%

- The average of both models matches their individual predictions.
Thus, both models are well-fitted to both train & test data.



Grid Search & Hyperparameters Analysis: Random Forest

Model	Hyperparamters Baseline	Hyperparameters Attempted
Random Forest	<code>n_estimators=100, max_depth=14, min_samples_leaf=10, max_features='sqrt'</code>	<code>n_estimators={64,100,200} max_features={2,3,4} bootstrap = {True,False} oob_score = {True,False}</code>

- **Goal:** try different number of trees, number of features used to consider best split, try utilizing the whole dataset to build tree
- **Results:**
 - with this grid search, could only achieve accuracy of 75%.
 - **It appears that `max_features` has the most effect on accuracy for this dataset. Utilizing the square root (features) returns best accuracy.**
 - Can achieve 99% accuracy if `max_depth` & `min_samples_leaf` are set to default. However, this takes too long computationally. There is tradeoff between the time and model accuracy.



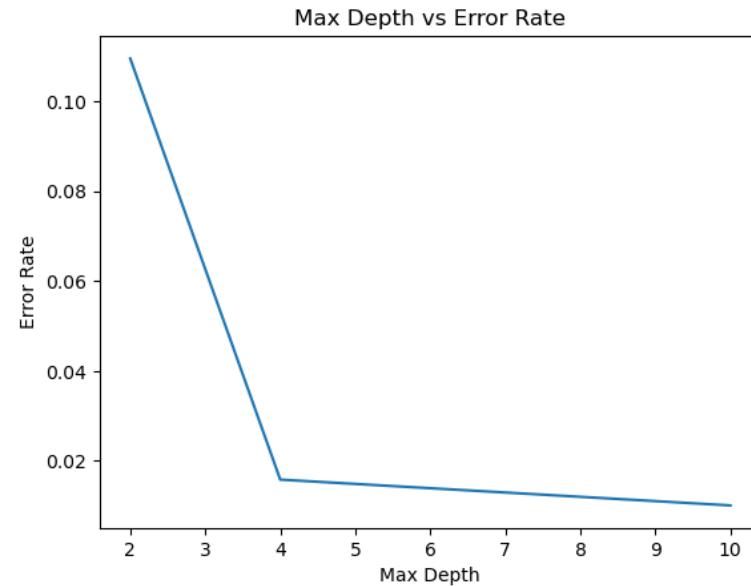
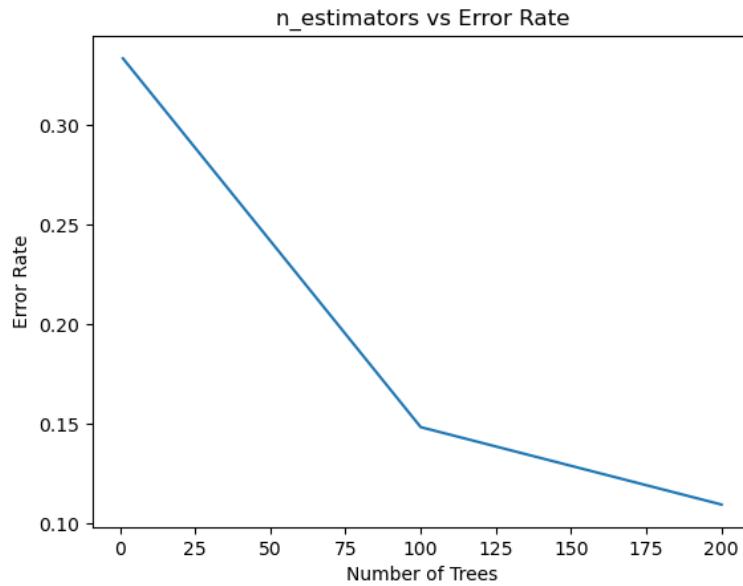
Grid Search & Hyperparameters Analysis: AdaBoost

Model	Hyperparamters Baseline	Hyperparameters Attempted
AdaBoost	<code>n_estimators=1,</code> <code>max_depth=2,</code> <code>learning_rate=1</code>	<code>n_estimators={1,100,200}</code> <code>max_depth={2,4,10}</code> <code>learning_rate = {0.5,1,2}</code>

- **Goal:** try different number of estimators for ensemble learning, different tree depths, and the weight (learning rate) applied to each classifier



Grid Search & Hyperparameters Analysis: AdaBoost



- Results: with the baseline hyperparameters, accuracy was ~75%. **200 estimators performed best, with accuracy at 89% for depth = 2. Accuracy continued to improve as depth increased.**



Future Work

- Try SVM model to compare with random forest & AdaBoost
 - This will provide us with 4 models of comparison: **logistic reg, random forest, AdaBoost, and SVM**
- Is there any further feature engineering that we can do to the dataset? Will it improve accuracy?



Conclusion



Random forest and AdaBoost baseline models return the best accuracy & fit. AdaBoost performs slightly better.



In order to achieve such high accuracy, hyperparameters for random forest & AdaBoost were tuned.



Will create SVM model for further comparison.



Determine best overall model and further fine tune as able.



Links

- Jupyter Notebook Codes:

<https://github.com/molls1889/DataMiningCheckpoint2>

- Video Link:

<https://github.com/molls1889/DataMiningCheckpoint2>